

# Provable Safety in Augmentation: A Production-Grade Framework of Neurorights, Lyapunov Control, and Rust-Based Enforcement

## Theoretical Foundations of Neural-Roping and Safety Envelopes

The development of a robust safety framework for human augmentation hinges on establishing a set of provably safe operating conditions that protect the operator, referred to as the "organic\_cpu," from unsafe access by external AI models . This section details the theoretical underpinnings of this framework, focusing on the BCI/EEG/MCI domain as the foundational entry point for protection . The approach integrates advanced mathematical concepts from control theory and formal logic to create dynamic, adaptive, and verifiable safety envelopes, moving beyond static thresholds to a more sophisticated model of biophysical and cognitive integrity. The central paradigm is one of "neural-rope," which emphasizes a safe, governed connection rather than invasive actuation, directly addressing concerns about mental privacy and integrity [10](#) [11](#) .

A cornerstone of this theoretical framework is the concept of the **corridor predicate**, an extension of the base admissibility predicate  $A_{H,C}$  over the fundamental bio-corridor polytope  $(E, M_{prot}, S_{bio}, \theta, \Delta T)$ , where  $E$  represents energy budget,  $M_{prot}$  protein stress markers,  $S_{bio}$  systemic inflammation,  $\theta$  thermal load, and  $\Delta T$  time duration . For the BCI/EEG domain, this polytope is extended with a new, domain-specific axis representing electrophysiological activity, denoted as  $\Phi_{EEG}$  . This axis captures the power spectral density of brain signals within specific frequency bands, such as NREM gamma or frontal theta, which are known to correlate with cognitive load and task engagement [50](#) [55](#) . The creation of a new ALN particle, such as `bio.corridor.eeg.gamma.v1`, formalizes this extended corridor, defining explicit bounds for  $\Phi_{EEG}$  and linking them to measurable physiological proxies and rollback triggers . This process grounds abstract safety requirements in concrete, observable quantities, a critical step for transitioning from high-level rights declarations to legally and technically enforceable contracts [1](#) .

Within this extended corridor, several key theoretical constructs are employed to ensure safety. First are the **Natural-Boundary Inequalities**. These are mathematical expressions that define the biophysical "do-not-cross" surfaces of the safety envelope . They represent the limits of the body's ability to compensate for metabolic and thermal stress induced by neural stimulation or processing load. For instance, an inequality might state that the product of EEG-derived duty cycle and local thermal proxy must remain below a threshold derived from studies on heat dissipation in neural tissue <sup>58</sup> . These inequalities are not arbitrary; they are anchored to empirical evidence through a bundle of ten hex tags that link each bound to established literature or internal data on CMRO (cerebral metabolic rate of oxygen), IL-6 cytokine levels, or other relevant biomarkers <sup>17 104</sup> . This rigorous evidence anchoring ensures that the theoretical boundaries have a basis in real-world physiology, enhancing the credibility and legal defensibility of the safety claims .

Second, the framework introduces **Lyapunov-style duty laws** to govern the dynamic adjustment of operational parameters. Instead of relying on simple binary checks ("if above limit, shut down"), this approach uses a continuous, feedback-driven control law. The proposed update rule,  $u_{k+1} = u_k - \eta \frac{\partial V}{\partial u}$ , adjusts the "duty"  $u$  (which could represent signal processing frequency, decision-making latency, or another workload metric) based on the gradient of a Lyapunov function  $V(u)$  . A common choice for  $V(u)$  is the squared deviation from a safe duty level,  $V(u) = (u - u_{safe})^2$  . This formulation ensures that when the system's state deviates from the safe region, the control law actively drives it back towards equilibrium, providing asymptotic stability <sup>38</sup> . This principle is analogous to Lyapunov optimization techniques used in resource management for edge computing, where the goal is to maintain system stability while meeting performance objectives <sup>43</sup> . By implementing this in the `lyapunov_descent()` method of the guard crates, the system achieves a form of self-regulation that is more nuanced and resilient than simple threshold-based alarms .

Third, the concept of **Viability Kernels** is introduced to manage the long-term health of the organic\_cpu, particularly concerning cognitive and metabolic load . A viability kernel is a set of initial states from which there exists at least one control strategy that can keep the system's trajectory within a predefined "safe" set indefinitely <sup>38</sup> . In this context, the "safe set" corresponds to biophysically acceptable ranges for energy consumption, protein stress, inflammation, and thermal load. The viability kernel defines the initial operating conditions and allowable transitions that respect these long-term constraints. Calculating this kernel provides a geometric characterization of all possible "safe" operational paths, allowing the system to proactively select trajectories that avoid irreversible harm. This is a significant theoretical advancement, framing augmentation not merely as a short-term technical problem but as a long-horizon biological challenge that requires careful

management of cumulative stress <sup>17</sup>. The integration of viability theory with controlled-invariant set geometry offers a powerful tool for proving the long-horizon safety of state transitions between different modes of operation, such as rest, training, and elite-sport performance .

Finally, these mathematical constructs are encoded into **ALN-encoded Corridor Predicates** that serve as formal contracts for neurorights . An ALN shard, such as `bio.corridor.eeg.gamma.v1.aln`, is a structured document that specifies the bounds, units, and logical relationships within the safety corridor . It includes not only the numerical limits for axes like EEG band power but also embedded clauses and flags, such as prohibiting the raw capture of narrative thoughts to protect mental privacy <sup>17</sup>. These ALN particles become the source of truth for all enforcement mechanisms, from compile-time checks in Rust to runtime decisions by the Virta-Sys orchestrator. The use of a formal grammar for interface telemetry and rejection rules, compiled into controllers and policy checkers, is a direct response to the need for more specific and practically applicable legal frameworks for neurotechnology <sup>1 17</sup>. This transforms abstract rights into concrete, machine-verifiable constraints.

Theoretical Construct	Mathematical Form / Description	Role in Safety Framework
Corridor Predicate	Extension of base polytope $(E, M_{prot}, S_{bio}, \theta, \Delta T)$ with domain-specific axes like $\Phi_{EEG}$ .	Defines the multi-dimensional safe operating space for a given augmentation task.
Natural-Boundary Inequalities	Biophysical constraints (e.g., relating EEG duty to thermal load) anchored to empirical data .	Establishes hard, provable "do-not-cross" limits for safety, grounded in physiology.
Lyapunov-Style Duty Law	Feedback control law $u_{k+1} = u_k - \eta \frac{\partial V}{\partial u}$ using a Lyapunov function $V(u)$ .	Provides dynamic, self-correcting regulation of operational load to maintain stability.
Viability Kernel	Set of initial states from which a control strategy can keep the system within a safe region indefinitely <sup>38</sup> .	Ensures long-term, sustainable operation by respecting cumulative biophysical stress limits.
ALN-Encoded Predicates	Formal shards (e.g., <code>bio.corridor.eeg.gamma.v1.aln</code> ) containing bounds, units, evidence tags, and neurorights clauses .	Creates a verifiable, auditable, and machine-executable contract for safety and rights.

This combination of theoretical tools creates a multi-layered defense. Natural boundaries provide the ultimate physical limits, viability kernels ensure long-term sustainability, and Lyapunov laws provide the active control needed to navigate complex operational scenarios. All are codified in ALN, creating a complete, traceable, and enforceable safety doctrine for neurorights-governed neural augmentation. The success of this framework depends critically on the quality of the evidence anchoring the natural boundary inequalities, a topic explored further in subsequent sections.

# Practical Implementation via Rust Guard Crates and Telemetry

The theoretical advancements in safety envelopes and control laws must be translated into immediately deployable, production-grade software artifacts to be effective. This section outlines the practical implementation strategy, centered on a family of domain-specific **Rust guard crates** that enforce the corridor predicates and duty laws at both compile-time and during runtime orchestration . This approach leverages the unique features of the Rust programming language—such as its ownership model, strong typing, and compile-time verification—to build systems with verifiable guarantees of safety and security [23](#) [24](#) . The implementation is designed to produce tangible outputs, including Prometheus metrics for monitoring and audit-ready logs for accountability, ensuring the framework is not just a theoretical exercise but a functional component of the augmented cybernetic ecosystem.

The core of the practical implementation is the **Rust Guard Crate**, a library designed to encapsulate the logic for a specific type of safety constraint. For the BCI/EEG domain, this would be a crate named `bio_eeg_guard`. Each crate exposes a set of typed structs and traits that mirror the structure of the corresponding ALN specification . For example, the ALN particle `bio.corridor.eeg.gamma.v1` would inform the design of a `EegCorridorBundle` struct, which aggregates all necessary telemetry data points such as slices of power spectral density, accumulated duty cycles, and a thermal proxy derived from metabolic rate estimates . This tight mapping ensures that any change to the safety corridor's definition in the ALN is automatically reflected in the software architecture.

The primary interface for these crates is a trait, such as `EegGuardKernel`, which defines the essential methods for safety enforcement . This trait serves as a clear API for the higher-level Virta-Sys orchestrator to interact with the guards. The key methods include:

- `fn admissible(&self, env, bundle) -> bool;` : This function takes the current environmental context and the telemetry bundle, then evaluates whether the operation is permitted according to the corridor predicate. It implements the logical inequalities defined in the ALN shard, returning ‘true’ only if all constraints, including those related to psych-risk and EEG duty, are satisfied .
- `fn lyapunov\_descent(&self, state, telemetry) -> State;` : This method implements the Lyapunov-style duty law. When the system detects that it is approaching or has entered a state of high risk, this function calculates a new, safer duty level or operational state that guarantees a descent towards the stable, safe region of the

state space . This provides the dynamic, corrective behavior required for robust safety.

- `fn chat\_knowledge\_factor(&self, evidence: EvidenceBundle) -> f32;` : This method computes a scalar "knowledge factor" that quantifies the confidence in the safety assessment. It rewards configurations with strong evidence bundles, low bioimpact, positive eco outcomes, and compliance signals, providing a metric for the quality and reliability of the operational context .

By defining these behaviors as a trait, the system allows for multiple, interchangeable implementations of the safety logic, facilitating testing and future algorithmic improvements. The use of Rust's type system allows for the creation of specialized wrapper types, such as the proposed `NeurorightsBound<T>` and `RiskEnvelope<T>`, which embed the safety constraints directly into the type signature of variables. This enables the compiler itself to catch many potential violations at compile-time, a powerful mechanism for preventing errors before deployment [23](#) .

To support this enforcement layer, a comprehensive **telemetry and metrics** infrastructure is essential. The system is designed to export Prometheus metrics that provide real-time visibility into the health and safety status of the augmentation protocols . For the BCI/EEG domain, this includes counters and gauges for metrics like `eeg_guard_duty`, `eeg_guard_reject_total`, and `eeg_guard_rollback_total`. Similarly, the nanoswarm crate would expose metrics like `nanoswarm_guard_admissible` and `nanoswarm_guard_thermal_margin` . These metrics are crucial for several reasons: 1. **Operational Monitoring:** They allow operators and automated systems to continuously monitor compliance with safety corridors and detect incipient problems. 2. **Performance Analysis:** They provide data for analyzing the effectiveness of the Lyapunov control laws and identifying patterns of constraint violation. 3. **Audit and Forensics:** They form part of the audit-ready logs, providing a detailed, timestamped record of all safety-related events that can be used for debugging, incident analysis, and regulatory reporting .

The daily development loop mandates that for every new theoretical predicate created in an ALN shard, exactly one corresponding element is implemented in the Rust crate: a new field in the telemetry bundle, a new inequality in the `admissible()` method, and a new term in the `lyapunov_descent()` function . This strict pairing ensures that the theoretical model and its software implementation remain perfectly synchronized. Furthermore, rigorous testing is a non-negotiable part of the process. Property tests are used to verify that the `admissible()` function correctly rejects operations when telemetry values exceed the defined bounds, while snapshot tests ensure that the

Lyapunov controller converges to the safe duty level `u_safe` from various starting points, demonstrating stable behavior .

This practical implementation strategy is designed for immediate utility. The resulting Rust crates can be directly integrated into the existing neurorights firewall and the Virta-Sys orchestration layer, hardening stakeholder protection across all biophysical-blockchain routes and enabling secure cross-system collaboration . The emphasis on verifiable outputs like Prometheus metrics and audit-ready logs makes the system transparent and accountable, a critical requirement for building trust in advanced augmentation technologies. The table below summarizes the relationship between the theoretical constructs and their practical counterparts.

Theoretical Construct	Corresponding Rust Implementation	Purpose & Outcome
Corridor Predicate ( <code>bio.corridor.eeg.gamma.v1.aln</code> )	<code>EegCorridorBundle</code> struct and <code>admissible()</code> trait method.	Aggregates telemetry data and evaluates compliance with the safety corridor at runtime.
Natural-Boundary Inequalities	Inequalities implemented within the <code>admissible()</code> method.	Provides hard-coded, non-negotiable safety checks against biophysical limits.
Lyapunov-Style Duty Law	<code>lyapunov_descent()</code> trait method.	Dynamically adjusts operational parameters to ensure stability and drive the system toward a safe state.
ALN-Encoded Predicates	Neurorights clauses within the ALN shard and corresponding logic in the Rust guard.	Enforces formal contracts for safety and neurorights, making them machine-executable.
Prometheus Metrics	Counters and gauges exported by the guard crate (e.g., <code>eeg_guard_reject_total</code> ).	Enables real-time monitoring, performance analysis, and auditing of safety compliance.
Audit-Ready Logs	Logging statements triggered by rejections, rollbacks, and duty adjustments.	Creates a detailed, immutable record of all safety-related events for accountability.

This systematic approach, combining formal specification with rigorously tested, production-grade code, bridges the gap between theoretical safety science and practical engineering, delivering a framework that is not only theoretically sound but also ready for deployment and verification.

# The CyberRank-KO Governance Spine: A Unified Ecosystem Conscience

While the BCI/EEG domain serves as the frontline for protection, the overarching governance of the entire augmented ecosystem requires a more generalized, domain-agnostic framework. This is provided by the **CyberRank vector integrated with Knowledge Object (KO) subnets**, which functions as a sovereign governance spine . Its primary purpose is to establish a unified ranking system across multiple dimensions—including safety, legal robustness, biomechanical evidence, psych-risk, and eco impact—that can be queried by all four rotating domains (BCI/EEG, Nanoswarm, Neuromorphic, Smart-City) . This spine acts as a cross-system conscience, ensuring that actions and policies in one domain adhere to the same global standards of safety and neurorights protection as those in another, thereby preventing fragmentation and bias.

The core of this governance model is the **CyberRank vector**, a low-dimensional vector that assigns a rank score to each KO within its subnet along several key axes . This is analogous to Google's PageRank algorithm but adapted for a knowledge graph where nodes are not web pages but particles, policies, or augmentations, and edges are weighted by factors like safety, evidence quality, and legal robustness . The computation involves a damped random walk over the graph of KOs, where the damping factor (e.g., 0.85) determines the probability of continuing the walk versus teleporting to a random node, ensuring convergence to a stable probability distribution . This approach is aligned with emerging standards for AI governance, such as those being developed by ISO/IEC JTC 1/SC 42, which emphasize transparency, explainability, and the establishment of trustworthy AI systems [87](#) [88](#) [107](#).

To manage complexity and scale, the system is organized into **KO subnets**, which are local directed graphs of KOs relevant to a specific domain or context . For instance, there would be a **BCI -KO -Subnet**, a **Nanoswarm -KO -Subnet**, and so on. The CyberRank engine computes a separate, normalized CyberRank vector for each subnet, reflecting the internal structure and relative importance of KOs within that specific context . This modular approach prevents the computational explosion that would come from trying to rank a single, monolithic global graph while still allowing for cross-subnet comparison through a shared semantic model.

The key architectural pattern for integrating the rotating domains with this spine is a **projection and querying mechanism**. Each domain does not need to implement its own

ranking algorithm. Instead, it defines a projection layer that maps its own specific metrics into the shared CyberRank vector fields and queries the spine for rankings . For example:

- A BCI job's safety profile (corresponding to `r\_bio` in the CyberRank vector) is projected from its telemetry and corridor predicate compliance.
- A swarm-node's audit coverage and policy adherence (`r\_obs`) are mapped into the observability dimension of the vector.
- The overall Risk-of-Harm (RoH  $\leq 0.3$ ) ceiling is enforced globally, regardless of the domain .

This pattern ensures a consistent, unified language for value and safety assessment across disparate systems. A KO in the BCI subnet and a KO in the smart-city subnet can be compared on a common scale because they are both evaluated against the same five axes of the CyberRank vector. This is crucial for enabling secure asset movement and collaboration between systems, as routing and authorship decisions can be made based on a holistic, cross-domain evaluation of safety and rights .

A proof-of-concept implementation of this model is provided in the form of a `crates/cyberrank_ko` Rust library . This crate contains three main modules: 1. **model.rs**: Defines the core data structures, including `KoId` for identifiers, `CyberRankEntry` for individual rank scores, `CyberRankVector` for the output vector, and `KoEdge` for the weighted connections between nodes . 2. **engine.rs**: Contains the core logic for computing the rank vector. The `PowerIterationEngine` struct implements the `CyberRankEngine` trait and executes the power iteration algorithm to solve the PageRank equation . It handles normalization, convergence checks using L1 distance, and the special case of dangling nodes (nodes with no outgoing edges) . 3. **lib.rs**: Exposes the public API and includes unit tests to verify correctness, such as confirming that the computed rank vector always sums to 1.0 and that the rankings respect the structural properties of simple test graphs .

While this provides a solid foundation, the model faces several challenges. The **computational complexity** of the power iteration method, with its `max_iterations` and `tolerance` parameters, must be carefully managed for very large, dynamic graphs representing city-scale networks . Numerical stability is also a concern, especially with ill-conditioned graphs. Furthermore, the **weighting of the CyberRank vector axes** is a critical design choice. If the weights are not calibrated to reflect societal values accurately, the system could introduce subtle biases, favoring certain augmentations or domains over others and potentially undermining the goal of equitable neurorights protection. Finally, there is a risk of **semantic drift**, where the meaning of the vector components ("safety," "legal robustness") could evolve differently in each domain over

time. To mitigate this, the CyberRank vector schema itself should be formally defined and version-controlled in an ALN shard to ensure consistency.

Despite these challenges, the CyberRank-KO spine offers a powerful solution for scalable, decentralized governance. By providing a stable, domain-agnostic ranking layer, it allows the four rotating domains to innovate independently while ensuring that their innovations are ultimately subject to the same overarching principles of safety and rights. It decouples the application logic from the governance logic, a design principle that enhances modularity and resilience. This architecture positions the CyberRank vector as the ecosystem's conscience, a persistent and computable representation of its collective commitment to protecting the augmented citizen.

## Paired Deliverables: ALN Shards and Audit-Ready Artifacts

The core innovation of this research program lies in its insistence on producing **paired deliverables**: for every new theoretical safety construct, a corresponding, immediately usable software artifact must be created alongside it. This tightly coupled workflow ensures that abstract mathematical ideas are rapidly translated into verifiable, production-grade components, bridging the chasm between academic research and industrial implementation. The two primary forms of these paired deliverables are the **Augmented Logic Notation (ALN) shards** for formal specification and the **Rust guard crates** for executable enforcement, complemented by a suite of **audit-ready artifacts** including Prometheus metrics and detailed logs.

The ALN shard is the formal specification document for a new safety corridor or predicate. Following the daily development loop, a new branch is created (e.g., `feat/2026-10-26-bci-corridor-v2`), and a new ALN file is committed (e.g., `specs/alm/bio/bio.corridor.eeg.gamma.v2.aln`). This file is not just a list of numbers; it is a rich, evidence-backed artifact. It begins by extending the base admissibility predicate and adding a new domain-specific axis, such as  $\Phi_{EEG}$  for EEG band power. Crucially, for each new bound defined on this axis, the specification must be accompanied by a bundle of **ten hex evidence tags** <sup>17</sup>. These tags, formatted as `0x?????????`, are cryptographic hashes that anchor the theoretical bound to a specific piece of empirical evidence—a research paper, a clinical trial dataset, or an internal validation study <sup>17</sup>. This practice directly addresses the call for more specific, scientifically-grounded formulations of neurorights, transforming subjective safety claims

into objective, falsifiable hypotheses <sup>1</sup>. The ALN shard also includes ALN contracts like `biosafeguard!` and `privacyscope!` to encode neurorights clauses, such as prohibitions on capturing raw narrative EEG, directly into the formal specification <sup>17</sup>.

Simultaneously, the Rust guard crate is updated. Using the new ALN shard as a blueprint, a developer implements exactly one new feature. This could be adding a new field to the telemetry bundle struct (e.g., `pub eeg_gamma_density: f32`), adding a corresponding inequality check in the `admissible()` method, and introducing a new term into the `lyapunov_descent()` function. This disciplined, one-to-one correspondence between the ALN and the Rust code ensures that the implementation never gets ahead of the specification and vice-versa. The resulting guard crate becomes a reusable, domain-specific library that can be plugged into Virta-Sys or any other orchestrator that needs to enforce BCI safety constraints. The use of Rust provides an additional layer of assurance, as its compile-time checks can enforce invariants specified in the ALN, such as ensuring that a `NeurorightsBound<T>` is always properly initialized <sup>23</sup>.

Beyond the core ALN and Rust files, the daily loop mandates the generation of a third set of paired deliverables: **audit-ready artifacts**. These are essential for accountability, debugging, and regulatory compliance. The first category is **Prometheus metrics**, which are standardized time-series databases for monitoring system health. For the BCI loop, this includes counters for `eeg_guard_reject_total` (number of rejected operations) and `eeg_guard_rollback_total` (number of times the duty was auto-degraded), and gauges for `eeg_guard_duty` (current operational load). These metrics are exposed by the guard crate and can be scraped by a central monitoring system, providing a real-time pulse on the system's safety posture. They are designed to be exportable into dashboards for both swarm-nodes and KO subnets, offering a unified view of system health across domains.

The second category of audit-ready artifacts is **detailed, timestamped logs**. Every significant safety-related event should trigger a log entry. This includes successful operations, rejections due to corridor violations, rollbacks initiated by the Lyapunov controller, and changes to the operational duty level. These logs must be structured and machine-readable, likely in JSON format, to facilitate automated analysis and forensics. They serve as an immutable record of the system's behavior, which is invaluable for investigating incidents, demonstrating compliance with neurorights charters, and improving the underlying safety models over time. The generation of these logs is a key component of producing a "hardened" system that protects stakeholders across all biophysical-blockchain routes.

Finally, the entire day's contribution is summarized and stamped in a `research-data-manifest.json` file. This manifest records the new ALN corridor, the corresponding knowledge-factor score, and the hex-stamp, creating a permanent, cryptographically verifiable record of the research progress . This manifests the abstract concept of a "knowledge-factor" into a concrete, measurable outcome. The knowledge-factor, a scalar  $F \in (0,1)$ , is computed based on several components, including the novelty of the contribution, the quality of the supporting evidence, the improvement in safety, and alignment with neurorights and eco-impact goals <sup>17</sup>. For example, a run that successfully maps magnetic duty and temperature into a new corridor inequality with clear rollback triggers might receive a high score like `0.93` . This scoring mechanism incentivizes high-quality, impactful research by making it visible and quantifiable. The table below illustrates the typical paired deliverables produced in a single day's cycle.

Artifact Type	Example File/Component	Content and Purpose
Formal Specification	<code>bio.corridor.eeg.gamma.v2.aln</code>	Defines new corridor bounds for EEG gamma band power, with 10 hex evidence tags and neurorights clauses .
Executable Code	<code>bio_eeg_guard crate</code>	Implements the <code>admissible()</code> and <code>lyapunov_descent()</code> methods based on the ALN specification .
Monitoring Metrics	Prometheus metrics ( <code>eeg_guard_*</code> )	Counters and gauges for tracking rejections, rollbacks, and current duty levels for real-time health monitoring .
Audit Trail	Structured log entries	Timestamped records of all safety-related events (rejections, rollbacks, etc.) for accountability and forensics .
Progress Manifest	<code>research-data-manifest.json</code>	A log file that registers the day's contributions, including the hex-stamp and computed knowledge-factor .

This systematic, paired approach ensures that the research program is not only innovative but also pragmatic and verifiable. It produces a growing corpus of formal specifications and production-ready code, all meticulously documented and monitored, creating a robust and transparent foundation for neurorights-governed augmentation.

## Strategic Synthesis and Future Research Trajectory

This research program presents a comprehensive and deeply integrated strategy for developing a safe and ethically governed infrastructure for human augmentation. Its strength lies in the synergistic interplay between three core pillars: the **BCI/EEG/MCI augmentation loop** as the foundational domain for protection; the **CyberRank-KO governance spine** as a sovereign, domain-agnostic arbiter; and the **integrated implementation stack** of ALN specifications and Rust guard crates that brings theory to

life. This synthesis reveals a clear pathway from abstract rights to verifiable, machine-enforceable code, addressing the critical need for specificity and practical applicability in the nascent field of neurorights <sup>1</sup>.

The most significant strategic insight is the project's attempt to operationalize neurorights. Rather than relying solely on high-level constitutional amendments, such as Chile's pioneering inclusion of neurorights in its constitution <sup>102103</sup>, this framework embeds these rights directly into the software and governance layers of the cybernetic ecosystem. The `NeurorightsBound<T>` and `biosafeguard!` macros are examples of rights-as-code, where protections like mental privacy are transformed into enforceable software contracts <sup>17</sup>. The CyberRank vector, with its axis for "legal robustness," provides a mechanism for computationally evaluating and prioritizing these rights across a distributed network of autonomous agents. This represents a profound shift from declarative rights to procedural, verifiable ones.

The methodology itself is a masterclass in systematic innovation. The rotating daily loop provides a structured, iterative process for continuous improvement, cycling through different domains (BCI, nanoswarm, neuromorphic, smart-city) while maintaining deep focus on one per day. The strict requirement for paired deliverables—producing an ALN shard and a corresponding Rust guard module simultaneously—creates a powerful feedback loop. It forces researchers to think constantly about the practical implications of their theoretical work and ensures that abstractions are always grounded in implementable reality. This approach mitigates the risk of creating purely academic theories that are disconnected from engineering needs. The introduction of a daily "knowledge-factor" score further institutionalizes this focus on novelty and utility, creating a culture of measurable progress.

However, the framework is not without its inherent challenges and risks, which primarily lie in the epistemological and social domains. The entire system rests on a complete and accurate model of the human body and mind. The success of the natural-boundary inequalities and viability kernels depends entirely on the quality of the evidence anchoring them. Any gaps, inaccuracies, or biases in the underlying physiological and psychological data could lead to catastrophic failures, where the system either imposes unnecessary restrictions or fails to prevent genuine harm. The project's heavy reliance on evidence anchoring highlights this dependency; without a robust, continuously updated, and vetted database linking telemetry to biological and psychological outcomes, the system's validity is questionable <sup>76 83</sup>.

Furthermore, the proposed CyberRank-KO governance spine, while elegant in its design, carries the risk of creating a powerful, centralized authority. While intended to be

domain-agnostic, the algorithm and its parameters (weights, damping factor) could inadvertently introduce systemic biases, leading to the preferential treatment of certain types of augmentations or users over others. Ensuring the fairness, transparency, and resistance to manipulation of this spine will be paramount. The project's alignment with international efforts like the UNESCO Recommendation on AI Ethics is a positive step towards establishing a normative foundation, but it does not eliminate the technical and political challenges of governing a complex, emergent system [116117](#).

Looking forward, the project has laid out a clear path for future exploration through its ten identified frontier topics . These topics extend the current methodology in several key directions: 1. **Cross-Domain Coupling:** Topics like "Neuro-eco CHAT invariants for nanoswarm-city coupling" and "Multi-corridor CHAT-stabilized negotiation" aim to create deeper, more meaningful interactions between the rotating domains, moving beyond simple telemetry sharing to true system-of-systems optimization . 2. **Advanced Control and Learning:** Research into "Evidence-aware Lyapunov gain scheduling" and "Biophysical renormalization of corridors under continual learning" seeks to make the safety controls themselves adaptive, capable of becoming more aggressive as confidence in the underlying models grows . 3. **Formal Verification and Governance:** Topics such as "Neuro-CHAT thermodynamics" and "Cross-jurisdiction diffmap algebra" push into more formal territory, aiming to create provable guarantees of safety and to build robust, verifiable pipelines for research and upgrades that are sovereign and non-seizable .

In conclusion, the research goal provides a visionary yet meticulously detailed roadmap for building a trustworthy augmentation ecosystem. It successfully marries the rigor of control theory with the pragmatism of modern software engineering and the necessity of ethical governance. The immediate priority should be the deep implementation of the BCI/EEG loop, producing the first complete cycle of ALN shard and Rust guard crate as a proof-of-concept. Concurrently, the CyberRank-KO spine should be validated on a small, simulated subnet. Addressing the epistemological risks by building a world-class evidence base and the social risks by designing a fair and transparent governance algorithm will be the ultimate determinants of the framework's success.

---

## Reference

1. The protection of mental privacy in the area of neuroscience [https://www.europarl.europa.eu/RegData/etudes/STUD/2024/757807/EPRS\\_STU\(2024\)757807\\_EN.pdf](https://www.europarl.europa.eu/RegData/etudes/STUD/2024/757807/EPRS_STU(2024)757807_EN.pdf)
2. (PDF) Artificial Intelligence and Neuroscience [https://www.researchgate.net/publication/388087540\\_Artificial\\_Intelligence\\_and\\_Neuroscience\\_Transformative\\_Synergies\\_in\\_Brain\\_Research\\_and\\_Clinical\\_Applications](https://www.researchgate.net/publication/388087540_Artificial_Intelligence_and_Neuroscience_Transformative_Synergies_in_Brain_Research_and_Clinical_Applications)
3. The state of clinical trials of implantable brain–computer ... [https://www.researchgate.net/publication/384317855\\_The\\_state\\_of\\_clinical\\_trials\\_ofImplantable\\_brain-computer\\_interfaces](https://www.researchgate.net/publication/384317855_The_state_of_clinical_trials_ofImplantable_brain-computer_interfaces)
4. American Journal of Physical Medicine & Rehabilitation [https://journals.lww.com/ajpmr/fulltext/2014/03001/abstracts\\_of\\_scientific\\_papers\\_and\\_posters.1.aspx](https://journals.lww.com/ajpmr/fulltext/2014/03001/abstracts_of_scientific_papers_and_posters.1.aspx)
5. Brain Sci., Volume 15, Issue 6 (June 2025) – 127 articles <https://www.mdpi.com/2076-3425/15/6>
6. Abstracts from the 2025 Annual Meeting of the American ... <https://journals.sagepub.com/doi/full/10.1177/15459683251360052>
7. 2019 ACVIM Forum Research Abstract Program - PMC <https://PMC6766498/>
8. 2019 AAPM&R Annual Assembly Abstracts <https://onlinelibrary.wiley.com/doi/10.1002/pmrj.12271>
9. Medicina, Volume 61, Issue 5 (May 2025) – 173 articles <https://www.mdpi.com/1648-9144/61/5>
10. Protecting Brain Privacy in the Age of Neurotechnology [https://www.researchgate.net/publication/384971350\\_Protecting\\_Brain\\_Privacy\\_in\\_the\\_Age\\_of\\_Neurotechnology\\_Policy\\_Respo\\_nses\\_and\\_Remaining\\_Challenges](https://www.researchgate.net/publication/384971350_Protecting_Brain_Privacy_in_the_Age_of_Neurotechnology_Policy_Respo_nses_and_Remaining_Challenges)
11. What Should We Do With People Who Cannot or ... [https://www.researchgate.net/publication/353686682\\_What\\_Should\\_We\\_Do\\_With\\_People\\_Who\\_Cannot\\_or\\_Do\\_Not\\_Want\\_to\\_Be\\_Protected\\_From\\_Neurotechnological\\_Threats/fulltext/610a8daf169a1a0103db64d9/What-Should-We-Do-With-People-Who-Cannot-or-Do-Not-Want-to-Be-Protected-From-Neurotechnological-Threats.pdf](https://www.researchgate.net/publication/353686682_What_Should_We_Do_With_People_Who_Cannot_or_Do_Not_Want_to_Be_Protected_From_Neurotechnological_Threats/fulltext/610a8daf169a1a0103db64d9/What-Should-We-Do-With-People-Who-Cannot-or-Do-Not-Want-to-Be-Protected-From-Neurotechnological-Threats.pdf)

12. On pharmacological neuroenhancement as part of the new ... [https://www.researchgate.net/journal/Frontiers-in-Psychology-1664-1078/publication/372463985\\_On\\_pharmacological\\_neuroenhancement\\_as\\_part\\_of\\_the\\_new\\_neurorights\\_pioneering\\_legislation\\_in\\_Chile\\_a\\_perspective/links/6628303543f8df018d2549ba/On-pharmacological-neuroenhancement-as-part-of-the-new-neurorights-pioneering-legislation-in-Chile-a-perspective.pdf](https://www.researchgate.net/journal/Frontiers-in-Psychology-1664-1078/publication/372463985_On_pharmacological_neuroenhancement_as_part_of_the_new_neurorights_pioneering_legislation_in_Chile_a_perspective/links/6628303543f8df018d2549ba/On-pharmacological-neuroenhancement-as-part-of-the-new-neurorights-pioneering-legislation-in-Chile-a-perspective.pdf)
13. The Impact of Artificial Intelligence on Human Thought <https://arxiv.org/pdf/2508.16628.pdf>
14. (PDF) Freedom of Thought in the Age of Neuroscience [https://www.researchgate.net/publication/261950057\\_Freedom\\_of\\_Thought\\_in\\_the\\_Age\\_of\\_Neuroscience](https://www.researchgate.net/publication/261950057_Freedom_of_Thought_in_the_Age_of_Neuroscience)
15. (PDF) Brain-computer interfaces and education: the state of ... [https://www.researchgate.net/publication/335486095\\_Brain-computer\\_interfaces\\_and\\_education\\_the\\_state\\_of\\_technology\\_and\\_imperatives\\_for\\_the\\_future](https://www.researchgate.net/publication/335486095_Brain-computer_interfaces_and_education_the_state_of_technology_and_imperatives_for_the_future)
16. Cognitive Enhancement: Methods, Ethics, Regulatory ... [https://www.researchgate.net/publication/26308471\\_Cognitive\\_Enhancement\\_Methods\\_Ethics\\_Regulatory\\_Challenges](https://www.researchgate.net/publication/26308471_Cognitive_Enhancement_Methods_Ethics_Regulatory_Challenges)
17. (PDF) Unveiling the Neurotechnology Landscape [https://www.researchgate.net/publication/372353690\\_Unveiling\\_the\\_Neurotechnology\\_Landscape\\_Scientific\\_Advancements\\_Innovations\\_and\\_Major\\_Trends](https://www.researchgate.net/publication/372353690_Unveiling_the_Neurotechnology_Landscape_Scientific_Advancements_Innovations_and_Major_Trends)
18. Nancy JECKER | Professor (Full) | PhD | Research profile <https://www.researchgate.net/profile/Nancy-Jecker-2>
19. Recommendations for Responsible Development and ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC8081770/>
20. Artificial Intelligence in Science and Society: The Vision of ... <https://ieeexplore.ieee.org/iel8/6287639/10820123/10839366.pdf>
21. Technology Vision 2024 <https://www.accenture.com/content/dam/accenture/final/accenture-com/document-2/Accenture-Tech-Vision-2024.pdf>
22. 1 - Amazon S3 <https://s3.eu-central-1.amazonaws.com/deepset.ai-farm-models/0.4.1/glove-english-uncased-6B/vocab.txt>
23. Machine Learning-Based Vulnerability Detection in Rust ... <https://www.mdpi.com/2504-4990/7/3/79>
24. microflow: an efficient rust-based inference engine <https://arxiv.org/pdf/2409.19432.pdf>
25. A Low-Latency Optimization of a Rust-Based Secure ... <https://www.mdpi.com/1424-8220/22/22/8700>

26. Use of 3D Printing Techniques to Fabricate Implantable ... <https://pubs.acs.org/doi/10.1021/acsmmeasurescian.3c00028>
27. Saad Motahhir Badre Bossoufi Editors | PDF <https://www.scribd.com/document/640638207/UUntitled>
28. words\_SG\_up to 2020.txt [https://zenodo.org/record/5516252/files/words\\_SG\\_up to 2020.txt](https://zenodo.org/record/5516252/files/words_SG_up to 2020.txt)
29. cmnt\_vocab.txt [https://www.cs.cmu.edu/~ark/blog-data/data/blog\\_data\\_v1\\_0/dk/hbc\\_data/data/cmnt\\_vocab.txt](https://www.cs.cmu.edu/~ark/blog-data/data/blog_data_v1_0/dk/hbc_data/data/cmnt_vocab.txt)
30. Toward cross-subject and cross-session generalization in ... <https://www.sciencedirect.com/science/article/pii/S0925231224011251>
31. Cross-Subject Motor Imagery Electroencephalogram ... <https://www.mdpi.com/2306-5354/12/5/495>
32. Past, Present, and Future of EEG-Based BCI Applications - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC9101004/>
33. Neurotechnologies for Brain-Machine Interfacing <https://standards.ieee.org/wp-content/uploads/import/documents/presentations/ieee-neurotech-for-bmi-standards-roadmap.pdf>
34. Integrating Target Domain Convex Hull with MMD for Cross ... <https://www.mdpi.com/2078-2489/17/1/15>
35. Brainwave Biometrics: A Secure and Scalable Brain- ... <https://www.mdpi.com/2673-2688/6/9/205>
36. Extreme manufacturing of minimally invasive intracortical ... <https://iopscience.iop.org/article/10.1088/2631-7990/ae2579>
37. What Should We Do With People Who Cannot or Do Not ... <https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2021.703092/full>
38. Deep Reinforcement Learning-Driven Robust Adaptive ... <https://ieeexplore.ieee.org/iel8/6287639/11323511/11314514.pdf>
39. Deep Reinforcement Learning for Power Converter Control <https://ieeexplore.ieee.org/iel8/8782709/10803118/11197642.pdf>
40. Validation of Sliding Mode and Passivity Control in High- ... <https://ieeexplore.ieee.org/iel7/6287639/10380310/10347200.pdf>
41. Impact-Resilient Orchestrated Robust Controller for Heavy- ... <https://ieeexplore.ieee.org/iel8/3516/11317928/10993327.pdf>
42. Neural Networks in a Softcomputing Framework <https://link.springer.com/content/pdf/10.1007/1-84628-303-5.pdf>

43. Cost optimization in edge computing: a survey - Springer Link <https://link.springer.com/article/10.1007/s10462-024-10947-4>
44. From MPC-Based to End-to-End (E2E) Learning-Based Control ... <https://ieeexplore.ieee.org/iel7/6287639/9668973/09771222.pdf>
45. 34th Annual Computational Neuroscience Meeting: CNS\*2025 <https://link.springer.com/article/10.1007/s10827-025-00915-4>
46. Weighting Factors Design of Model Predictive Control for ... <https://ieeexplore.ieee.org/iel8/6287639/10380310/10583912.pdf>
47. Adaptive Control for Simultaneous Tracking and Stabilization ... <https://link.springer.com/content/pdf/10.1007/s10846-023-01908-0.pdf>
48. Hyperkinetic Movement Disorders - 2023 <https://movementdisorders.onlinelibrary.wiley.com/doi/10.1002/mds.29542>
49. Poster Presentations - 2015 - Wiley Online Library <https://movementdisorders.onlinelibrary.wiley.com/doi/10.1002/mds.26295>
50. Current implications of EEG and fNIRS as functional ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC11629311/>
51. EEG classification for neurological disorders using ... <https://www.nature.com/articles/s41598-025-33535-0>
52. Advancing EEG-based brain-computer interface ... <https://www.sciencedirect.com/science/article/pii/S2590238524002480>
53. Concurrent fNIRS and EEG for Brain Function Investigation <https://www.mdpi.com/1424-8220/22/15/5865>
54. Resting-state EEG associated with clinical measures to ... <https://www.frontiersin.org/journals/neurology/articles/10.3389/fneur.2025.1577393/pdf>
55. EEG provides insights into motor control and neuroplasticity ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC11421965/>
56. Brain computer interfaces for cognitive enhancement in older ... <https://link.springer.com/article/10.1186/s12877-025-05676-4>
57. Brain-to-Text Decoding: A Non-invasive Approach via Typing <https://arxiv.org/pdf/2502.17480.pdf>
58. Structural and Functional Designs of Advanced Neural ... <https://pubs.acs.org/doi/10.1021/acsami.5c15407>
59. Review on EEG – Based Authentication Technology - Zhang <https://onlinelibrary.wiley.com/doi/10.1155/2021/5229576>
60. (PDF) Progress in Brain Computer Interface: Challenges ... [https://www.researchgate.net/publication/349597769\\_Progress\\_in\\_Brain\\_Computer\\_Interface\\_Challenges\\_and\\_Opportunities](https://www.researchgate.net/publication/349597769_Progress_in_Brain_Computer_Interface_Challenges_and_Opportunities)

61. Arxiv今日论文 | 2025-12-11 [http://lonepatient.top/2025/12/11/arxiv\\_papers\\_2025-12-11](http://lonepatient.top/2025/12/11/arxiv_papers_2025-12-11)
62. Computer Science <https://arxiv.org/list/cs/new>
63. Biotechnology and Human Enhancement - Springer Link <https://link.springer.com/content/pdf/10.1007/978-94-024-2296-2.pdf>
64. Pattern Classification (2nd Edition): | Guide books <https://dl.acm.org/doi/book/10.5555/954544>
65. Planet Mozilla <https://planet.mozilla.org/>
66. glove.6B.100d.txt-vocab.txt <https://worksheets.codalab.org/rest/bundles/0xadf98bb30a99476ab56ebff3e462d4fa/contents/blob/glove.6B.100d.txt-vocab.txt>
67. Abstracts from The 39th Annual Symposium of the National ... <https://journals.sagepub.com/doi/10.1089/neu.2022.29126.abstracts>
68. Abstracts of Scientific Papers and Posters Presented at the ... [https://journals.lww.com/jisprm/fulltext/2020/03002/abstracts\\_of\\_scientific\\_papers\\_and\\_posters.1.aspx](https://journals.lww.com/jisprm/fulltext/2020/03002/abstracts_of_scientific_papers_and_posters.1.aspx)
69. Untitled | PDF | Medicine | Educational Technology <https://www.scribd.com/document/636488929/Untitled>
70. Adaptive neuromodulation dialogues - PubMed Central - NIH <https://pmc.ncbi.nlm.nih.gov/articles/PMC12715846/>
71. Multimodal behavioral data predict stroke patient's ... <https://pubmed.ncbi.nlm.nih.gov/41134943/>
72. Interim Safety Profile From the Feasibility Study of the ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC10074470/>
73. The 14th International Joint Conference on Natural ... <https://aclanthology.org/events/ijcnlp-2025/>
74. Machine Learning Feb 2025 <https://www.arxiv.org/list/cs.LG/2025-02?skip=200&show=1000>
75. Table of Contents <https://ieeexplore.ieee.org/iel7/9577055/9577056/09578286.pdf>
76. Strategies to Improve the Robustness and Generalizability ... <https://www.mdpi.com/2673-7426/5/2/20>
77. 机器学习2025\_6\_6 <http://www.arxivdaily.com/thread/68259>
78. Use of BCI Systems in the Analysis of EEG Signals for ... <https://dl.acm.org/doi/full/10.1145/3757732>
79. Review article EEG-based emotion recognition systems <https://www.sciencedirect.com/science/article/pii/S2405844024075169>

80. Brain-Computer Interface: Advancement and Challenges - PMC <https://PMC8433803>
81. Brain-Computer Interfaces <https://iopscience.iop.org/book/edit/978-0-7503-5914-6/chapter/bk978-0-7503-5914-6ch4.pdf>
82. Current Status, Challenges, and Possible Solutions of EEG ... <https://www.frontiersin.org/journals/neurorobotics/articles/10.3389/fnbot.2020.00025/full>
83. institutional analysis of consumer neurotechnology regulation [https://www.researchgate.net/publication/388517713\\_INSTITUTIONAL\\_ANALYSIS\\_OF\\_CONSUMER\\_NEUROTECHNOLOGY\\_REGULATION\\_MEDIATING\\_THIRD-PARTY\\_AUTONOMY](https://www.researchgate.net/publication/388517713_INSTITUTIONAL_ANALYSIS_OF_CONSUMER_NEUROTECHNOLOGY_REGULATION_MEDIATING_THIRD-PARTY_AUTONOMY)
84. Towards Usable Electroencephalography-based Brain- ... [https://inria.hal.science/tel-01416980/file/HDR\\_FabienLotte\\_Final.pdf](https://inria.hal.science/tel-01416980/file/HDR_FabienLotte_Final.pdf)
85. In the realm of the hybrid brain: Human Brain and AI <https://arxiv.org/pdf/2210.01461>
86. Current Trends, Challenges, and Future Research ... <https://www.mdpi.com/2414-4088/7/10/95>
87. ISO/IEC JTC 1/SC 42 - Artificial intelligence <https://www.iso.org/committee/6794475.html>
88. 国际标准化机构联合发布人工智能管理系统国际标准 [https://www.ccpitbj.org/web/static/articles/catalog\\_40fcc0367c729df20181af782c2307f1/article\\_ff8080818abfb00a018e1b97b95f0d42/ff8080818abfb00a018e1b97b95f0d42.html](https://www.ccpitbj.org/web/static/articles/catalog_40fcc0367c729df20181af782c2307f1/article_ff8080818abfb00a018e1b97b95f0d42/ff8080818abfb00a018e1b97b95f0d42.html)
89. Data analytics subgroup report - EMA [https://www.ema.europa.eu/en/documents/report/data-analytics-subgroup-report\\_en.pdf](https://www.ema.europa.eu/en/documents/report/data-analytics-subgroup-report_en.pdf)
90. 国家标准项目 <https://std.samr.gov.cn/gb/search/gbDetailed?id=91B7130A0EEDF892E05397BE0A0A978E>
91. Artificial intelligence - Standards by ISO/IEC JTC 1/SC 42 <https://www.iso.org/committee/6794475/x/catalogue/p/%200/u/1/w/0/d/0>
92. Digital Twin AI: Opportunities and Challenges from Large ... <https://arxiv.org/html/2601.01321v1>
93. Computational approaches to Explainable Artificial ... <https://www.sciencedirect.com/science/article/pii/S1566253523002610>
94. (PDF) Long-term unsupervised recalibration of cursor ... [https://www.researchgate.net/publication/398449459\\_Long-term\\_unsupervised\\_recalibration\\_of\\_cursor-based\\_intracortical\\_brain-computer\\_interfaces\\_using\\_a\\_hidden\\_Markov\\_model](https://www.researchgate.net/publication/398449459_Long-term_unsupervised_recalibration_of_cursor-based_intracortical_brain-computer_interfaces_using_a_hidden_Markov_model)
95. Closed-Loop Cognitive Stress Regulation Using Fuzzy ... <https://ieeexplore.ieee.org/iel7/6287639/9312710/09492148.pdf>

96. CRISPR and Artificial Intelligence in Neuroregeneration <https://www.mdpi.com/1422-0067/26/19/9409>
97. Transforming Sleep Monitoring: Review of Wearable and ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC11853583/>
98. Biomedicines, Volume 14, Issue 1 (January 2026) <https://www.mdpi.com/2227-9059/14/1>
99. Invited Faculty Abstracts from the International ... <https://onlinelibrary.wiley.com/doi/full/10.1111/ner.12958>
100. UC San Diego Electronic Theses and Dissertations <https://escholarship.org/content/qt17r640dh/qt17r640dh.pdf>
101. Soft Bioelectronics Based on Nanomaterials | Chemical Reviews <https://pubs.acs.org/doi/10.1021/acs.chemrev.1c00531>
102. Chile: pioneering the protection of neurorights [https://unesdoc.unesco.org/ark:/48223/pf0000380275\\_eng](https://unesdoc.unesco.org/ark:/48223/pf0000380275_eng)
103. Neurorights in the Constitution: from neurotechnology to ethics ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC11491849/>
104. Chilean Supreme Court ruling on the protection of brain ... <https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2024.1330439/full>
105. ISO/IEC JTC 1/SC 42人工智能分技术委员会第一次全会在京 ... <http://lti.ruc.edu.cn/sy/xwdt/48a5666aa8f4479797d8cf43395c3fa8.htm>
106. SC 42 – Artificial Intelligence [https://www.itu.int/en/ITU-T/extcoop/ai-data-commons/Documents/ISO\\_IEC%20JTC1%20SC%2042%20Keynote\\_Wael%20Diab.pdf](https://www.itu.int/en/ITU-T/extcoop/ai-data-commons/Documents/ISO_IEC%20JTC1%20SC%2042%20Keynote_Wael%20Diab.pdf)
107. Artificial intelligence - ISO/IEC JTC 1/SC 42 <https://www.iso.org/committee/6794475/x/catalogue/p/1/u/0/w/0/d/0%7D>
108. Pilot results from continuous EEG-fNIRS study - PubMed Central <https://pmc.ncbi.nlm.nih.gov/articles/PMC8594770/>
109. Resting state EEG biomarkers of cognitive decline ... <https://pmc.ncbi.nlm.nih.gov/articles/PMC7864432/>
110. Machine learning in biosignal analysis from wearable devices <https://pubs.rsc.org/en/content/articlehtml/2025/mh/d5mh00451a>
111. ACNP 58th Annual Meeting: Poster Session I - PMC <https://pmc.ncbi.nlm.nih.gov/articles/PMC6957924/>
112. 2025 CNS Meeting Abstract - 2025 - Annals of Neurology <https://onlinelibrary.wiley.com/doi/10.1002/ana.78038?af=R>
113. Neurocritical Care Society Abstracts from the 22nd Annual ... <https://link.springer.com/content/pdf/10.1007/s12028-024-02129-5.pdf>

114. Mobile Networks and Management - Springer Link <https://link.springer.com/content/pdf/10.1007/978-3-030-94763-7.pdf>
115. Artificial intelligence - ISO/IEC JTC 1/SC 42 <https://www.iso.org/es/contents/data/committee/67/94/6794475.html?view=catalogue>
116. Recommendation on the Ethics of Artificial Intelligence <https://unesdoc.unesco.org/ark:/48223/pf0000380455>
117. UNESCO Recommendation on the Ethics of Artificial ... <https://www.ohchr.org/sites/default/files/2022-03/UNESCO.pdf>
118. Recommendation on the Ethics of Artificial Intelligence <https://www.unesco.org/en/articles/recommendation-ethics-artificial-intelligence>
119. UNESCO Recommendation on AI Ethics <https://academy.evalcommunity.com/unesco-recommendation-on-ai-ethics/>
120. Draft text of the Recommendation on the Ethics of Artificial ... <https://unesdoc.unesco.org/ark:/48223/pf0000377898>
121. Draft text of the Recommendation on the Ethics of Artificial ... <https://unesdoc.unesco.org/ark:/48223/pf0000376713>
122. Brain-computer interfaces - ISO/IEC JTC 1/SC 43 <https://www.iso.org/committee/9082407.html>
123. Big Data Reference Architecture (ISO/IEC JTC 1 SC42) <https://www.linkedin.com/pulse/big-data-reference-architecture-isoiec-jtc-1-sc42-ray>
124. news\_vocab\_sorted.txt [https://www.cs.cmu.edu/afs/cs.cmu.edu/academic/class/15122-s14/www/14-interfaces/news\\_vocab\\_sorted.txt](https://www.cs.cmu.edu/afs/cs.cmu.edu/academic/class/15122-s14/www/14-interfaces/news_vocab_sorted.txt)
125. 3617623 12|3594859 time|3570693 year|3516017 publisher [http://svn.apache.org/repos/asf/lucene/dev/tags/realtime\\_DWPT\\_final\\_2011-05-02/solr/src/test-files/Top50KWiki.utf8](http://svn.apache.org/repos/asf/lucene/dev/tags/realtime_DWPT_final_2011-05-02/solr/src/test-files/Top50KWiki.utf8)